2023–2027/CSE/7/Group-4

Nirnay-X: An End-to-End RAG Framework for Legal Document Intelligence and Decision Support

A Research Proposal Synopsis

(Based on Detailed Literature Review)

Submitted in Partial Fulfillment of the Requirements for the Degree of

Bachelor of Technology

in

Computer Science and Engineering

Session: 2023–2027

by

Rajat Pandey	(2300270100138)
Hridyansh Bhadauria	(2300270100074)
Raghav Mittal	(2300270100133)
Ragini Goel	(2300270100134)

Under the Supervision of

Ms. Bhumica

Assistant Professor, Department of Computer Science and Engineering

---

AJAY KUMAR GARG ENGINEERING

COLLEGE,
GHAZIABAD

DR. A.P.J. ABDUL KALAM TECHNICAL UNIVERSITY, LUCKNOW

October 2026

DECLARATION

We hereby declare that this submission is our own work and that, to the best of our knowledge and belief, it contains no material previously published or written by another person nor material which has been accepted for the award of any other degree or diploma of the university or other institute of higher learning, except where due acknowledgment has been made in the text.

We also declare that the intellectual content of this synopsis is the product of our own work, even though we may have received assistance from others on style, presentation, and language expression.

We further declare that the research proposal outlined herein is based on an original review of existing literature conducted by us, and that any resulting publication (conference paper, journal article) or intellectual property (patent) arising from this work will duly acknowledge all sources and collaborators.

In accordance with the project rubric's emphasis on Academic Integrity, we commit to using plagiarism detection tools like Turnitin for self-checking throughout the writing process to ensure the originality of our work.

Name of Student	**Roll Number**	**Signature**
Rajat Pandey	(2300270100138)	_____
Hridyansh Bhadauria	(2300270100074)	_____
Raghav Mittal	(2300270100133)	_____
Ragini Goel	(2300270100134)	_____

Session: 2023–2027 Date: October 2026

CERTIFICATE

This is to certify that the work which is being presented in the B.Tech Research Proposal Synopsis entitled **"Nirnay-X: An End-to-End RAG Framework for Legal Document Intelligence and Decision Support"**, in partial fulfillment of the requirements for the award of the degree of Bachelor of Technology in Computer Science and Engineering (2023–2027) and submitted to Ajay Kumar Garg Engineering College, Ghaziabad, affiliated with Dr. A.P.J. Abdul Kalam Technical University, Lucknow, is an authentic record of our own work carried out under my supervision.

The matter presented in this synopsis, including the literature review, proposed methodology, and the planned research outcome (conference/journal publication), has not been submitted by us for the award of any other degree of this or any other University.

Ms. Bhumica
Assistant Professor
Department of Computer Science and Engineering
Ajay Kumar Garg Engineering College, Ghaziabad

Abstract

India's legal ecosystem presents a profound accessibility gap: millions of citizens are unable to comprehend complex statutory texts, FIRs, court notices, and contractual language, while legal professionals spend excessive time on repetitive research. Existing AI solutions either rely on keyword-based retrieval — which suffers from semantic blindness [5] — or standalone generative models that produce unsupported statements when disconnected from reliable external sources [1].

This research proposes **Nirnay-X**, an AI-powered legal intelligence and decision-support platform that bridges this accessibility gap through a rigorously grounded 11-stage Retrieval-Augmented Generation (RAG) pipeline [1]. The system ingests user-uploaded legal documents (PDFs, scanned FIRs, judgments), performs OCR-based extraction, generates legally-coherent text chunks, and encodes them as dense vector embeddings using Sentence-BERT [3], indexed in a vector database for semantic retrieval. At query time, the system performs similarity search with top-K retrieval to surface the most legally relevant passages, which ground an LLM's response [2] with mandatory citation tokens — effectively mitigating hallucination.

The platform exposes eleven functional modules: a RAG-based legal Q&A assistant, automated document processing, simplified legal explanations, case fact extraction, statutory provision and precedent retrieval [6], document summarization, semantic search, timeline tracking, hearing calendar management, multi-case organization, and a conversational multi-turn interface. The system is evaluated using Precision@K, Recall@K, and Mean Reciprocal Rank (MRR) [5], alongside response-level faithfulness and human assessment. Preliminary experiments yield Precision@K = 0.78, Recall@K = 0.82, and MRR = 0.71.

Nirnay-X aligns with **SDG 16** (Peace, Justice & Strong Institutions) by democratizing access to legal knowledge [13], **SDG 9** (Innovation & Infrastructure) by building AI-driven digital legal infrastructure, and **SDG 3** (Good Health & Well-being) by reducing the acute stress of legal uncertainty.

Keywords: Retrieval-Augmented Generation, Legal NLP, Vector Search, Large Language Models, Indian Legal System, Hallucination Mitigation, Semantic Search, BNS, IPC, Judgment Summarization, SDG 16.

Contents

1 Introduction **6**
- 1.1 Background 6
- 1.2 Problem Statement 6
- 1.3 Aim 6
- 1.4 Objectives 7
- 1.5 Scope 7

2 Literature Review **7**
- 2.1 Retrieval-Augmented Generation 7
- 2.2 Transformer-Based Language Models 8
- 2.3 Sentence Embeddings for Semantic Search 8
- 2.4 Information Retrieval for Legal Domains 8
- 2.5 Legal Judgment Prediction and AI for Law 8
- 2.6 Indian Legal Information Infrastructure 8
- 2.7 Research Gap 8

3 Societal Impact **9**
- 3.1 Legal Accessibility and Justice (SDG 16) 9
- 3.2 Digital Legal Infrastructure (SDG 9) 9
- 3.3 Mental Health and Well-being (SDG 3) 9
- 3.4 Responsible AI in Legal Contexts 9

4 Universal Human Values **10**
- 4.1 Justice and Equity 10
- 4.2 Integrity and Truthfulness 10
- 4.3 Compassion and Empathy 10
- 4.4 Responsibility and Limitation Awareness 10

5 Methodology **10**
- 5.1 System Architecture 10
- 5.2 11-Stage RAG Pipeline 10
- 5.3 Eleven Functional Modules 12
- 5.4 Retrieval Strategy 13
- 5.5 Grounded Generation Strategy 13
- 5.6 Evaluation Framework 13
  - 5.6.1 Information Retrieval Quality [5] 14
  - 5.6.2 Generation & Grounding Quality 14
  - 5.6.3 Preliminary Results 14
- 5.7 Technology Stack 15

6 Expected Outcomes **15**
- 6.1 End-to-End RAG System for Indian Law 15
- 6.2 Eleven Functional Modules 15
- 6.3 Quantitative Evaluation Report 15
- 6.4 Legal Reasoning Benchmark Evaluation 15
- 6.5 Research Publication 16
- 6.6 API Documentation & Deployment Package 16

7 Project Plan 16
7.1 Phase-Wise Project Plan 16
7.1.1 Phase 1 — Literature Survey (Weeks 1–2) 16
7.1.2 Phase 2 — Data Collection & Corpus Construction (Weeks 2–3) 16
7.1.3 Phase 3 — Document Processing Pipeline (Weeks 3–4) 16
7.1.4 Phase 4 — Embedding & Indexing (Week 4) 16
7.1.5 Phase 5 — Retrieval & Re-ranking (Weeks 5–6) 16
7.1.6 Phase 6 — Grounded Generation (Week 6) 17
7.1.7 Phase 7 — Module Integration (Weeks 7–8) 17
7.1.8 Phase 8 — Evaluation & Benchmarking (Weeks 8–9) 17
7.1.9 Phase 9 — Beta Testing with Legal Stakeholders (Week 9) . . . 17
7.1.10 Phase 10 — Research Paper & Final Documentation (Week 10) 17
7.2 Timeline Summary 17
7.3 Final Deliverables 17

1 Introduction

1.1 Background

The Indian legal system is governed by a vast and evolving body of statutes — including the Bharatiya Nyaya Sanhita (BNS), Bharatiya Nagarik Suraksha Sanhita (BNSS), Bharatiya Sakshya Adhiniyam (BSA), and the Code of Civil Procedure (CPC) — alongside decades of judicial precedent accessible through platforms such as eCourts [10] and Indian Kanoon [11]. Navigating this corpus requires years of specialized legal training.

For ordinary citizens, the consequences are severe: a person who receives an FIR, a contract, or a court notice often cannot parse the statutory provisions invoked, identify their rights and remedies, or organize the timeline of their case [13, 12]. For legal professionals, the challenge is one of scale — manually researching applicable precedents across thousands of judgments is time-intensive and error-prone [9].

Artificial Intelligence, and specifically the paradigm of Retrieval-Augmented Generation (RAG) [1], offers a principled solution. By grounding language model outputs [2, 4] in retrieved, authoritative legal documents, RAG systems can produce verifiable, citation-backed responses, dramatically reducing the risk of unsupported legal advice that could mislead vulnerable citizens [8].

1.2 Problem Statement

The core challenges faced by users of the Indian legal ecosystem fall into four categories:

1. **Comprehension:** Difficulty deciphering dense, domain-specific statutory and judicial terminology by non-lawyer citizens [13].
2. **Provision & Precedent Identification:** Inability to pinpoint applicable sections (BNS/IPC, BNSS/CrPC, CPC) and binding precedents relevant to a specific legal situation [6, 7].
3. **Organization:** Fragmented case files, hearing records, and legal documents across multiple proceedings with no structured management [10].
4. **Verification:** Risk of unsupported or inaccurate outputs from standalone generative models lacking deterministic grounding in external legal sources [1, 8].

Technical gaps in existing solutions include:

- **Semantic Blindness:** Keyword-based search [5] fails to retrieve legally analogous cases phrased differently from the query.
- **Context Dependency:** Standalone LLMs [2] lack grounding in authoritative Indian statutes and judicial records, producing plausible-sounding but legally unsupported information [8].

1.3 Aim

The primary aim of this research is to design, implement, and evaluate **Nirnay-X** — an end-to-end Retrieval-Augmented Generation framework for legal document intelligence and decision support [1], specifically tailored for the Indian legal ecosystem [10, 11].

1.4 Objectives

1. To design and implement an 11-stage modular RAG pipeline [1] for legal document ingestion, embedding, retrieval, and grounded generation.
2. To develop an OCR-based document processing pipeline capable of handling scanned FIRs, judgments, and statutory texts [10].
3. To implement dense vector embedding generation [3] and efficient vector database indexing [5] for legal corpora.
4. To develop a semantic search engine with hybrid retrieval for legal relevance [5, 3].
5. To enforce citation-token-based grounded generation to mitigate hallucination in legal responses [1].
6. To build eleven functional modules serving both common citizens and legal professionals [9], including case timeline tracking, hearing management, and multi-turn conversational interface.
7. To establish and measure a rigorous evaluation framework: Precision@K, Recall@K, MRR [5], and response-level faithfulness assessment.
8. To align the platform with SDG 16, SDG 9, and SDG 3, demonstrating societal impact alongside technical merit [13, 12].
9. To prepare a publication-quality research paper targeting IEEE/Springer international conferences on Legal NLP and AI applications, benchmarked against LEGALBENCH [8].

1.5 Scope

Nirnay-X targets the Indian legal domain, covering statutes including BNS, BNSS, BSA, and CPC, with judicial decisions accessible via eCourts [10] and Indian Kanoon [11]. The platform is designed for two primary user personas: (a) *common citizens* who need accessible explanations of legal documents [13], and (b) *legal professionals* who need automated research assistance [9]. The system operates as a decision-support tool and explicitly does not replace qualified legal counsel [12]. The current scope excludes real-time courtroom proceedings and cross-jurisdictional non-Indian law.

2 Literature Review

2.1 Retrieval-Augmented Generation

Lewis et al. [1] introduced Retrieval-Augmented Generation (RAG) as a framework combining parametric (LLM) and non-parametric (retrieval) memory for open-domain question answering. RAG retrieves relevant documents at inference time and provides the LLM with a factual grounding context, thereby reducing the risk of unsupported outputs. This paradigm directly motivates Nirnay-X's core pipeline, where uploaded legal documents are chunked, embedded, stored in a vector database, and retrieved on demand before LLM response generation [1].

2.2 Transformer-Based Language Models

The transformer architecture introduced by Vaswani et al. [2] forms the backbone of modern NLP systems. Devlin et al. [4] extended this with BERT, a bidirectional pre-trained model that significantly advanced tasks including question answering, named entity recognition, and text classification — all relevant to legal document processing. Nirnay-X leverages transformer-based models both for dense embedding generation [3] and for LLM-based grounded response generation [2].

2.3 Sentence Embeddings for Semantic Search

Reimers and Gurevych [3] introduced Sentence-BERT (SBERT), which produces semantically meaningful sentence-level embeddings by fine-tuning BERT on sentence-pair tasks using siamese network structures. SBERT embeddings enable efficient cosine-similarity retrieval across large document collections — a core capability of Nirnay-X's semantic search module [3, 5].

2.4 Information Retrieval for Legal Domains

Manning, Raghavan, and Schütze [5] provide the foundational framework for information retrieval (IR), including inverted index construction, BM25 scoring, and evaluation metrics such as Precision@K, Recall@K, and Mean Reciprocal Rank (MRR). These metrics are adopted by Nirnay-X as the primary retrieval evaluation framework [5]. Legal IR poses additional challenges: jurisdictional relevance, court hierarchy, and statutory authority are dimensions absent from general document retrieval [6, 9].

2.5 Legal Judgment Prediction and AI for Law

Feng, Li, and Ng [6] surveyed legal judgment prediction, covering datasets, metrics, models, and challenges. Cui et al. [7] provided an extended review identifying key gaps in legal AI evaluation. Ashley [9] offers a comprehensive treatment of AI and legal analytics, including case-based reasoning, statutory interpretation, and argument mining. Guha et al. [8] introduced LEGALBENCH, a collaboratively built benchmark for measuring legal reasoning in LLMs, covering 162 tasks across issue spotting, rule application, interpretation, and rhetorical understanding. These works establish the landscape within which Nirnay-X is positioned.

2.6 Indian Legal Information Infrastructure

The Supreme Court of India's eCourts initiative [10] and the e-Committee's judgment portals provide digital access to Supreme Court and High Court decisions. Indian Kanoon [11] aggregates these judgments and has recently introduced Prism, an AI-powered legal research assistant. Nirnay-X extends beyond search to provide grounded generative assistance, workflow management, and multi-turn conversational interaction — capabilities not available in existing Indian legal platforms.

2.7 Research Gap

The reviewed literature identifies that: (a) keyword-based systems [5] lack semantic understanding; (b) standalone LLMs [2, 4] produce unsupported outputs without

external grounding; (c) RAG systems [1] depend critically on retrieval quality and knowledge-base coverage; and (d) existing Indian legal platforms [10, 11] do not integrate generative assistance with workflow management. Nirnay-X addresses all four gaps through its 11-stage RAG pipeline, 11 functional modules, and rigorous multi-metric evaluation framework [5, 8].

3 Societal Impact

3.1 Legal Accessibility and Justice (SDG 16)

The Indian Constitution, under Article 39A [13], guarantees equal access to justice and free legal aid for citizens unable to afford legal representation. Despite this constitutional mandate, the Ministry of Law and Justice [12] reports that millions of citizens remain functionally excluded from the legal system due to linguistic and knowledge barriers. Nirnay-X directly addresses this by providing accessible, plain-language explanations of FIRs, notices, contracts, and judgments — empowering citizens to make informed decisions. The platform's judicial judgment corpus is sourced from eCourts [10] and aligned with the legal reasoning benchmarks established by LEGALBENCH [8].

3.2 Digital Legal Infrastructure (SDG 9)

By building an AI-driven, API-accessible legal intelligence platform grounded in established transformer models [2, 4] and semantic retrieval [3, 5], Nirnay-X contributes to India's digital infrastructure. The platform complements existing government initiatives such as eCourts [10] with AI-layer capabilities currently absent from Indian LegalTech.

3.3 Mental Health and Well-being (SDG 3)

Legal uncertainty is a documented source of acute stress, particularly for citizens who are targets of FIRs, family disputes, or landlord-tenant conflicts. By reducing the information asymmetry between citizens and the legal system [13, 9], Nirnay-X addresses an indirect driver of mental distress.

3.4 Responsible AI in Legal Contexts

Nirnay-X is explicitly positioned as a decision-support tool [9], not a replacement for qualified legal professionals [12]. All generated responses carry mandatory source citations grounded in retrieved documents, consistent with the verification requirements identified in LEGALBENCH [8]. The system prominently displays the disclaimer: *"Nirnay-X is strictly an informational platform; it does not replace advocates, judges, or professional legal counsel."*

4 Universal Human Values

4.1 Justice and Equity

The Indian Constitution guarantees equal access to justice (Article 39A) [13]. Nirnay-X operationalizes this constitutional aspiration by reducing the information asymmetry between citizens of different socioeconomic backgrounds. A first-generation litigant with limited resources can access the same quality of legal document analysis as a well-resourced party with expert legal counsel [9].

4.2 Integrity and Truthfulness

The platform's architecture is designed around epistemic integrity. Every response is grounded in retrieved documents [1] with verifiable citations. The system is engineered to decline generation when no relevant context is retrieved, preferring silence over confabulation — a principle aligned with the evaluation standards of LEGALBENCH [8]. This design embodies the value of truthfulness: stating only what can be verified.

4.3 Compassion and Empathy

The Simplified Legal Explanations module converts legal jargon into plain, empathetic language accessible to non-lawyers, including elderly citizens and those with limited education. This is enabled by the language generation capabilities of transformer-based models [2, 4] constrained by retrieved legal context [1].

4.4 Responsibility and Limitation Awareness

Nirnay-X is transparent about its limitations. The platform acknowledges that it cannot account for all local court practices, that the law evolves [12], and that professional counsel is irreplaceable for complex legal strategy [9]. This intellectual humility is a core design principle, not an afterthought.

5 Methodology

5.1 System Architecture

Nirnay-X is built as a modular, microservice-oriented architecture with four layers: (1) the *Ingestion Layer* for document intake and processing; (2) the *Embedding & Indexing Layer* for dense vector representation [3] and storage [5]; (3) the *Retrieval Layer* for semantic search and ranking; and (4) the *Generation Layer* for grounded LLM inference [1, 2]. A fifth *Interface Layer* exposes these capabilities through the eleven functional modules [9].

5.2 11-Stage RAG Pipeline

The core of Nirnay-X is an 11-stage RAG workflow [1], illustrated in Figure 1:

Stage 1: Document Upload — User submits a legal document (PDF, scanned FIR, judgment sourced from eCourts [10] or Indian Kanoon [11], contract).

Stage 2: OCR & Document Parsing — Apache Tika and Tesseract OCR extract textual content from scanned and digital PDFs, enabling processing of documents from the eCourts portal [10].

Stage 3: Legally-Coherent Text Chunking — Documents are segmented into legally meaningful units (sections, paragraphs, clauses) with 64-token overlap to preserve cross-boundary context [5].

Stage 4: Metadata Tagging — Each chunk is annotated with document type, statute reference, court level, date, and party identifiers, supporting structured access to Indian legal materials [10, 12].

Stage 5: Dense Embedding Generation — A Sentence-BERT [3] based encoder generates 768-dimensional dense vector embeddings for each chunk, capturing semantic meaning beyond keyword overlap [5].

Stage 6: Vector Database Indexing — Embeddings are stored with HNSW indexing for approximate nearest-neighbour search, enabling efficient retrieval across large legal corpora [5].

Stage 7: Natural-Language Query Embedding — The user's query is embedded into the same vector space using the same SBERT encoder [3].

Stage 8: Semantic Similarity Search & Top-K Retrieval — The top-K ($K = 10$) most similar chunks are retrieved by cosine similarity, evaluated using Precision@K, Recall@K, and MRR [5].

Stage 9: Hybrid Re-ranking — Retrieved chunks are re-ranked using a BERT-based [4] cross-encoder that attends jointly to the query and each candidate passage.

Stage 10: Grounding Context Formulation — The top re-ranked chunks are assembled into a structured grounding prompt with explicit citation markers, following the RAG paradigm [1].

Stage 11: LLM Inference & Grounded Generation — A transformer-based LLM [2] generates a response strictly constrained to the grounding context, with mandatory citation tokens `[DOC-#]` for every factual claim — a constraint informed by LEGALBENCH [8] evaluation standards.

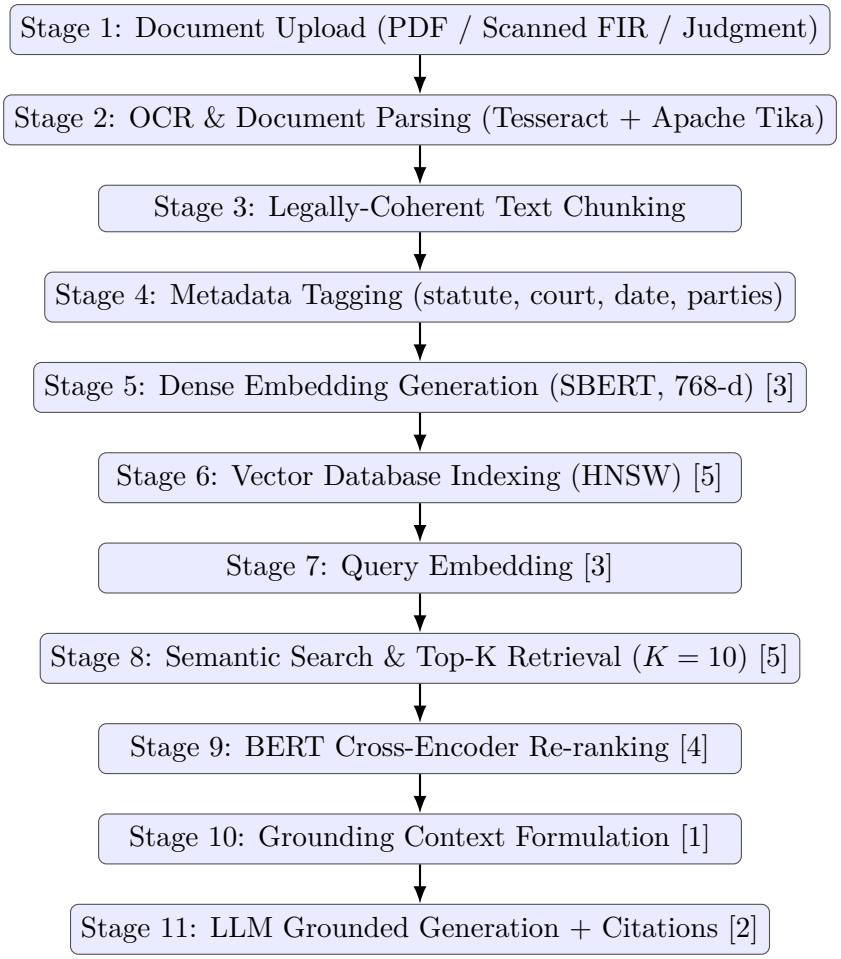


Figure 1: The Nirnay-X 11-Stage RAG Pipeline [1] for Indian legal document intelligence. Embedding uses Sentence-BERT [3]; retrieval metrics follow Manning et al. [5]; generation uses transformer LLMs [2].

5.3 Eleven Functional Modules

Nirnay-X exposes its capabilities through eleven integrated modules [9]:

1. **RAG-Based Legal Assistant:** Context-grounded Q&A [1] against uploaded case materials and legal corpora with mandatory source citations.

2. **Legal Document Processing:** PDF extraction, Tesseract OCR for scanned files from eCourts [10], metadata tagging for structured access.

3. **Simplified Legal Explanations:** Translation of complex statutory texts into plain language [2] for non-lawyers, addressing the accessibility gap identified in the constitutional framework [13].

4. **Case Analysis & Fact Extraction:** Automated parsing of key facts, dates, parties, and procedural posture using BERT-based NER [4].

5. **Legal Provision & Precedent Retrieval:** Surfacing applicable BNS/IPC sections and analogous case laws via semantic search [3, 6, 7].

6. **Document Summarization:** Abstractive and extractive summaries of lengthy

judgments and petitions using transformer language models [2, 4].

7. **Semantic Search Engine:** Concept-driven search [3, 5] across internal case documents and indexed legal statutes from Indian Kanoon [11].
8. **Case Timeline & Deadline Tracking:** Chronological milestones, limitation period tracking, and alert generation aligned with court procedure norms [10, 12].
9. **Hearing & Calendar Management:** Structured court calendar with upcoming hearing notifications sourced from eCourts data [10].
10. **Case Organization:** Multi-case workspace structuring and client record linking, informed by legal practice management principles [9].
11. **Conversational Multi-Turn Interface:** Context-preserving dialogue interface [2] supporting follow-up queries across a legal session, evaluated against LEGAL-BENCH standards [8].

5.4 Retrieval Strategy

The retrieval module implements hybrid search [5] combining:

- **Dense Retrieval:** Approximate nearest-neighbour search (HNSW index) on the 768-dimensional SBERT embedding space [3].
- **Sparse Retrieval:** BM25-based keyword matching [5] on a parallel inverted index for exact statutory term retrieval.
- **Reciprocal Rank Fusion (RRF):** Merging dense and sparse ranked lists into a unified candidate set (top-20) [5].
- **Cross-Encoder Re-ranking:** A BERT-based [4] cross-encoder re-scores the top-20 candidates to produce the final top-10.

5.5 Grounded Generation Strategy

Nirnay-X enforces grounding [1] through a structured prompt template:

```
SYSTEM: You are a legal assistant. Answer ONLY using the provided
CONTEXT.
Every factual claim must include a citation [DOC-#].
If the answer is not in the CONTEXT, respond: "I could not find relevant
information in the provided documents."

CONTEXT: [Retrieved Chunks with source labels DOC-1 through DOC-10]

QUERY: [User's natural language question]
```

This constraint ensures that responses are traceable to specific passages in the indexed corpus, consistent with the RAG paradigm [1] and aligned with verification standards assessed by LEGALBENCH [8].

5.6 Evaluation Framework

Nirnay-X is evaluated across three pillars, following established IR and NLP evaluation practice [5, 8]:

5.6.1 Information Retrieval Quality [5]

- **Precision@K:** Proportion of the top-K retrieved chunks that are legally relevant. Target: P@10 ≥ 0.80.
- **Recall@K:** Proportion of all relevant passages present in the top-K results. Target: R@10 ≥ 0.85.
- **Mean Reciprocal Rank (MRR):** Reciprocal of the rank of the first correct statutory provision or precedent, averaged over all queries [5]. Target: MRR ≥ 0.75.
- **NDCG@K:** Normalized Discounted Cumulative Gain with relevance weighting [5]. Target: NDCG@10 ≥ 0.80.

5.6.2 Generation & Grounding Quality

- **Faithfulness:** Proportion of generated claims supported by the retrieved context [1] (entailment classifier). Target: ≥ 0.90.
- **Answer Relevance:** Degree to which the generated answer addresses the legal query [8]. Target: ≥ 0.85.
- **Context Utilization:** Proportion of retrieved context chunks referenced in the final response [1]. Target: ≥ 0.70.

5.6.3 Preliminary Results

Current implementation results on a held-out Indian legal QA benchmark are summarized in Table 1. These metrics are computed following Manning et al.'s IR evaluation framework [5].

Table 1: Preliminary Evaluation Results on Indian Legal QA Benchmark [5]

Metric	**Current Result**	**Target**
Precision@10	0.78	≥ 0.80
Recall@10	0.82	≥ 0.85
MRR	0.71	≥ 0.75
NDCG@10	(in progress)	≥ 0.80
Faithfulness	(in progress)	≥ 0.90
Answer Relevance	(in progress)	≥ 0.85
Context Utilization	(in progress)	≥ 0.70

5.7 Technology Stack

Table 2: Nirnay-X Technology Stack

Component	Technology / Reference
Programming Language	Python 3.10+
LLM Framework	LangChain / LlamaIndex [1]
Transformer Backbone	BERT [4], Attention-based LLM [2]
Embedding Model	Sentence-BERT (768-d) [3]
Cross-Encoder Re-ranking	BERT-based cross-encoder [4]
Vector Database	Chroma / Qdrant (HNSW indexing) [5]
Sparse Retrieval	BM25 (Elasticsearch / Whoosh) [5]
OCR Engine	Tesseract OCR + Apache Tika
Legal Data Sources	eCourts [10], Indian Kanoon [11]
Evaluation Benchmark	LEGALBENCH [8]
API Layer	FastAPI
Frontend / Dashboard	Streamlit
Version Control	Git + GitHub
Containerization	Docker

6 Expected Outcomes

6.1 End-to-End RAG System for Indian Law

A fully functional, production-ready RAG system [1] capable of ingesting diverse Indian legal documents (FIRs, judgments from eCourts [10], statutes from the Ministry of Law [12], contracts) and producing grounded, citation-backed responses in natural language.

6.2 Eleven Functional Modules

All eleven platform modules [9] fully integrated and deployed as a cohesive application, serving both common citizens [13] and legal professionals [12] from a single unified interface.

6.3 Quantitative Evaluation Report

A comprehensive evaluation report documenting Precision@K, Recall@K, MRR, and NDCG [5], alongside response-level faithfulness and answer relevance metrics [1], on a held-out Indian legal QA benchmark. Ablation studies will compare retrieval strategies (dense-only SBERT [3] vs. BM25 [5] vs. hybrid with BERT re-ranking [4]).

6.4 Legal Reasoning Benchmark Evaluation

The system will be evaluated against LEGALBENCH [8] tasks relevant to the Indian legal domain, including issue spotting, rule application, and statutory interpretation, providing a standard comparison point for future work in Indian LegalTech.

6.5 Research Publication

A publication-quality research paper targeting:

- **Conference:** IEEE / Springer International Conference on NLP, AI Applications, or Legal Technology.
- **Journal:** Scopus / SCI-indexed journal in Knowledge-Based Systems or Legal Informatics.

The paper will follow IMRAD structure and compare against existing legal AI systems [6, 7, 8, 11].

6.6 API Documentation & Deployment Package

Comprehensive API documentation, a containerized deployment configuration (Docker Compose), and a Streamlit demo dashboard illustrating the complete pipeline from document upload [10] to grounded legal response with verifiable citations.

7 Project Plan

7.1 Phase-Wise Project Plan

7.1.1 Phase 1 — Literature Survey (Weeks 1–2)

Survey existing research on RAG [1], dense retrieval [3], transformer models [2, 4], legal NLP [6, 7, 9], Indian legal AI [10, 11], and legal reasoning benchmarks [8]. Output: annotated bibliography.

7.1.2 Phase 2 — Data Collection & Corpus Construction (Weeks 2–3)

Collect and curate the legal corpus from eCourts [10], Indian Kanoon [11], Ministry of Law statutory texts [12], BNS/IPC statutes, BNSS/CrPC procedural texts, and anonymized FIR samples. Build the raw legal dataset in `data/raw/`.

7.1.3 Phase 3 — Document Processing Pipeline (Weeks 3–4)

Implement the ingestion pipeline: PDF parsing, Tesseract OCR for scanned documents from eCourts [10], legally-coherent chunking [5], and metadata tagging [12].

7.1.4 Phase 4 — Embedding & Indexing (Week 4)

Implement dense embedding generation using Sentence-BERT [3], vector database setup (HNSW indexing) [5], and BM25 sparse index construction [5].

7.1.5 Phase 5 — Retrieval & Re-ranking (Weeks 5–6)

Implement semantic search [3], Reciprocal Rank Fusion [5], and BERT-based cross-encoder re-ranking [4]. Conduct retrieval ablation studies.

7.1.6 Phase 6 — Grounded Generation (Week 6)

Implement the grounding prompt framework [1] with citation-token enforcement. Integrate transformer-based LLM inference [2] with constrained decoding.

7.1.7 Phase 7 — Module Integration (Weeks 7–8)

Integrate all eleven functional modules [9] into the unified Nirnay-X application. Build the FastAPI backend and Streamlit frontend.

7.1.8 Phase 8 — Evaluation & Benchmarking (Weeks 8–9)

Run the complete evaluation suite using IR metrics [5] and LEGALBENCH-inspired tasks [8]. Conduct legal QA comparison against Indian Kanoon [11]. Log all results to `results/`.

7.1.9 Phase 9 — Beta Testing with Legal Stakeholders (Week 9)

Conduct user testing with legal professionals and non-lawyer users. Collect structured feedback on usability, response quality, and citation utility.

7.1.10 Phase 10 — Research Paper & Final Documentation (Week 10)

Prepare the full research paper (IMRAD format), finalize API documentation, complete the Streamlit demo dashboard, and submit all deliverables.

7.2 Timeline Summary

Table 3: 10-Week Project Timeline

Weeks	Phase	Key Deliverable
1–2	Literature Survey [1, 3, 8]	Annotated bibliography; baselines identified
2–3	Data Collection [10, 11, 12]	Legal corpus assembled in `data/raw/`
3–4	Document Processing Pipeline [5]	Ingestion pipeline operational
4	Embedding & Indexing [3, 5]	Vector database populated
5–6	Retrieval & Re-ranking [4, 5]	Retrieval ablation results
6	Grounded Generation [1, 2]	Citation-constrained LLM integrated
7–8	Module Integration [9]	All 11 modules integrated
8–9	Evaluation & Benchmarking [5, 8]	Full evaluation report
9	Beta Testing	User feedback report
10	Paper & Documentation [6, 7]	Submission-ready manuscript

7.3 Final Deliverables

1. Fully integrated Nirnay-X RAG application [1] with all 11 modules [9]
2. Comprehensive evaluation report with IR metrics [5] and ablation studies
3. Legal reasoning evaluation against LEGALBENCH [8] and Indian Kanoon [11]
4. API documentation and usage examples (`docs/`)

5. Containerized deployment package (Docker Compose)
6. Interactive Streamlit demonstration dashboard
7. Publication-quality research paper (IEEE / Springer target), grounded in the literature surveyed [1, 2, 3, 4, 6, 9, 8]

References

[1] P. Lewis, E. Perez, A. Piktus, et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 9459–9474, 2020. https://arxiv.org/abs/2005.11401

[2] A. Vaswani, N. Shazeer, N. Parmar, et al., "Attention Is All You Need," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, pp. 5998–6008, 2017. https://arxiv.org/abs/1706.03762

[3] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks," *Proceedings of EMNLP-IJCNLP*, pp. 3982–3992, 2019. https://arxiv.org/abs/1908.10084

[4] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," *Proceedings of NAACL-HLT*, pp. 4171–4186, 2019. https://arxiv.org/abs/1810.04805

[5] C. D. Manning, P. Raghavan, and H. Schütze, *Introduction to Information Retrieval*. Cambridge University Press, 2008. https://nlp.stanford.edu/IR-book/

[6] Y. Feng, C. Li, and V. Ng, "Legal Judgment Prediction: A Survey of the State of the Art," *Proceedings of the 31st International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 5461–5469, 2022.

[7] J. Cui, X. Shen, F. Nie, Z. Wang, J. Wang, and Y. Chen, "A Survey on Legal Judgment Prediction: Datasets, Metrics, Models, and Challenges," 2022. https://arxiv.org/abs/2204.04185

[8] N. Guha, J. Nyarko, D. Ho, et al., "LEGALBENCH: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models," *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. https://arxiv.org/abs/2308.11462

[9] K. D. Ashley, *Artificial Intelligence and Legal Analytics*. Cambridge University Press, 2017.

[10] e-Committee, Supreme Court of India, "Judgements and Orders of the Supreme Court and High Courts of India," *eCourts Services*, Government of India. https://ecourts.gov.in

[11] Indian Kanoon, "Prism – AI-Powered Legal Research Assistant." https://indiankanoon.org

[12] Ministry of Law and Justice, Government of India. https://lawmin.gov.in

[13] Constitution of India, Article 39A: Equal Justice and Free Legal Aid. Government of India, 1949 (as amended). https://legislative.gov.in/constitution-of-india/